


Surrogate calibration on optimized scenarios (weekly aggregates, n = 577)

Daganzo-LGB-Hybrid v2-aug (current model)

MAPE = 6.62% bias = -6.01% $R^2 = 0.9941$



LGB-logT v3 (previous generation)

MAPE = 8.32% bias = -8.02% $R^2 = 0.9544$

